

Assignment3

Cyril 20109216

20th of April 2020

Question a

We are interested in the following:

$$\begin{array}{ll} \arg \min_{\mathbf{f}(\mathbf{x})} & \mathbb{E}_{\mathbf{Y}|\mathbf{x}} \exp \left(-\frac{1}{K} (Y_1 f_1 + \dots + Y_K f_K) \right) \\ \text{subject to} & f_1 + \dots + f_K = 0. \end{array}$$

We can derive the following population minimizer and derive class probabilities as such:

The Lagrange of this constrained optimization problem can be written as:

$$\exp \left(-\frac{f_1(\mathbf{x})}{K-1} \right) \text{Prob}(c=1|\mathbf{x}) + \dots + \exp \left(-\frac{f_K(\mathbf{x})}{K-1} \right) \text{Prob}(c=K|\mathbf{x}) - \lambda (f_1(\mathbf{x}) + \dots + f_K(\mathbf{x})),$$

where λ is the Lagrange multiplier. Taking derivatives with respect to f_k and λ , we reach

$$\begin{array}{rcl} -\frac{1}{K-1} \exp \left(-\frac{f_1(\mathbf{x})}{K-1} \right) \text{Prob}(c=1|\mathbf{x}) - \lambda & = & 0, \\ & \vdots & \\ -\frac{1}{K-1} \exp \left(-\frac{f_K(\mathbf{x})}{K-1} \right) \text{Prob}(c=K|\mathbf{x}) - \lambda & = & 0, \\ f_1(\mathbf{x}) + \dots + f_K(\mathbf{x}) & = & 0. \end{array}$$

Solving this set of equations, we obtain the population minimizer

$$f_k^*(\mathbf{x}) = (K-1) \left(\log \text{Prob}(c=k|\mathbf{x}) - \frac{1}{K} \sum_{k'=1}^K \log \text{Prob}(c=k'|\mathbf{x}) \right), \quad k=1, \dots, K. \quad (4)$$

Thus,

$$\arg \max_k f_k^*(\mathbf{x}) = \arg \max_k \text{Prob}(c=k|\mathbf{x}),$$

which is the Bayes optimal classification rule in terms of minimizing the misclassification error. This justifies the use of this multi-class exponential loss function. Equation (4) also provides a way to recover the class probability $\text{Prob}(c=k|\mathbf{x})$ once $f_k^*(\mathbf{x})$'s are estimated, i.e.

$$\text{Prob}(c=k|\mathbf{x}) = \frac{e^{\frac{1}{K-1} f_k^*(\mathbf{x})}}{e^{\frac{1}{K-1} f_1^*(\mathbf{x})} + \dots + e^{\frac{1}{K-1} f_K^*(\mathbf{x})}}, \quad k=1, \dots, K.$$

Question b

Adaboost algorithm is the following:

Algorithm 1 *AdaBoost (Freund & Schapire 1997)*

1. Initialize the observation weights $w_i = 1/n$, $i = 1, 2, \dots, n$.

2. For $m = 1$ to M :

(a) Fit a classifier $T^{(m)}(\mathbf{x})$ to the training data using weights w_i .

(b) Compute

$$err^{(m)} = \sum_{i=1}^n w_i \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) / \sum_{i=1}^n w_i.$$

(c) Compute

$$\alpha^{(m)} = \log \frac{1 - err^{(m)}}{err^{(m)}}.$$

(d) Set

$$w_i \leftarrow w_i \cdot \exp \left(\alpha^{(m)} \cdot \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) \right), \quad i = 1, 2, \dots, n.$$

(e) Re-normalize w_i .

3. Output

$$C(\mathbf{x}) = \arg \max_k \sum_{m=1}^M \alpha^{(m)} \cdot \mathbb{I}(T^{(m)}(\mathbf{x}) = k).$$

Whereas the multiclass boosting algorithm is as follows:

Algorithm 2 *SAMME*

1. Initialize the observation weights $w_i = 1/n$, $i = 1, 2, \dots, n$.

2. For $m = 1$ to M :

(a) Fit a classifier $T^{(m)}(\mathbf{x})$ to the training data using weights w_i .

(b) Compute

$$err^{(m)} = \sum_{i=1}^n w_i \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) / \sum_{i=1}^n w_i.$$

(c) Compute

$$\alpha^{(m)} = \log \frac{1 - err^{(m)}}{err^{(m)}} + \log(K - 1). \quad (1)$$

(d) Set

$$w_i \leftarrow w_i \cdot \exp \left(\alpha^{(m)} \cdot \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) \right), \quad i = 1, \dots, n.$$

(e) Re-normalize w_i .

3. Output

$$C(\mathbf{x}) = \arg \max_k \sum_{m=1}^M \alpha^{(m)} \cdot \mathbb{I}(T^{(m)}(\mathbf{x}) = k).$$

Hence, we can conclude that the special case of the multiclass algorithm reduce to Adaboost as in Step (c),

the term $\log(K-1)$ where $K=2$ as in Adaboost vanishes.

Question c

```
error = function (w, misses)
{
  sum(w * misses)
}

alpha = function (err, K)
{
  if (err > 0.5) {
    1
  }
  else if (err <= 0) {
    20
  }
  else {
    log((1 - err)/err) + log(K - 1)
  }
}

weights = function (w, a, misses)
{
  tmp_w <- w * exp(a * misses)
  tmp_w/(sum(tmp_w))
}

samplePrediction = function (sample, A)
{
  votes.table <- votesTable(sample, A)
  prediction(votes.table)
}

weightedClassVote = function (k, sample, A)
{
  sum(A * (sample == k))
}

votesTable = function (sample, A)
{
  classes <- unique(sample)
  sapply(classes, function(k) weightedClassVote(k, sample,
    A))
}

prediction = function (votes.table)
{
  max.votes <- names(which(votes.table == max(votes.table)))
  ifelse(length(max.votes) == 1, max.votes, sample(max.votes,1))
}

finalPrediction = function (C, A)
{
  apply(C, 1, function(sample) samplePrediction(sample, A))
}
```

And finally the algorithm:

```
samme = function (formula, data, test, m = 5)
{
```

```

outcome.label <- toString(formula[[2]])
y <- data[, outcome.label]
C <- matrix(nrow = nrow(test), ncol = m)
n <- nrow(data)
w <- rep(1/n, n)
A <- numeric(m)
K <- nlevels(y)
for (i in 1:m) {
  t <- data[sample(n, n, replace = T, prob = w), ]
  t[, outcome.label] <- droplevels(t[, outcome.label])
  fit <- ranger(formula, data = t)
  C[, i] <- as.character(predict(fit, test)$predictions)
  h <- predict(fit, data)$predictions
  levels(h) <- levels(y)
  misses <- as.numeric(h != y)
  err <- error(w, misses)
  A[i] <- alpha(err, K)
  w <- weights(w, A[i], misses)
}
finalPrediction(C, A)
}

```

Load data

```

data(iris)
colnames(iris)[length(colnames(iris))] = 'y'
train_index <- sample(1:nrow(iris), nrow(iris)*0.75)
train = iris[train_index,]
test = iris[-train_index,]

```

Results

```


preds_same = samme(y ~., train, test)
model_gbm = gbm(y ~., data = train)

```

```
## Distribution not specified, assuming multinomial ...
```

```
best.iter <- gbm.perf(model_gbm, method = "OOB")
```

```
## OOB generally underestimates the optimal number of iterations although predictive performance is rea
```



```

preds_gbm = predict(model_gbm,newdata=test, n.trees = best.iter,type='link')
preds_gbm2 = matrix(preds_gbm,ncol=nlevels(train$y))
colnames(preds_gbm2) = colnames(preds_gbm[,1])
preds_gbm2 = apply(preds_gbm2,1,function(x) colnames(preds_gbm2)[which.max(x)])
acc_gbm = accuracy(preds_gbm2,test$y)
acc_same = accuracy(preds_same,test$y)
print(acc_gbm)

```

```
## [1] 0.9210526
```

```
print(acc_same)
```

```
## [1] 0.9210526
```

Hence, we can conclude that both algorithms have nearly similar performance.

Sources <https://web.stanford.edu/~hastie/Papers/samme.pdf> https://www.cs.ucr.edu/~eamonn/time_series_data/